

# F# Streams

A lightweight F#/C# library for efficient functional-style pipelines on streams of data.



# About Me

- Gian Ntzik (aka Jan Dzik)
- @anirothan
- Imperial College, Nessos

# About Nessos

- ISV based in Athens, Greece
- .NET experts
- Open source F# projects
  - {m}brace
  - FsPickler, Vagrant, and of course Streams

# Motivation

Make functional data query pipelines FAST

# LinqOptimizer

An automatic query optimizer-compiler for Sequential and Parallel LINQ.

<https://github.com/nessos/LinqOptimizer>

# LinqOptimizer

- compiles LINQ queries into fast loop-based imperative code
- speedups of up to 15x

# Example

## The query

```
var query = (from num in nums.AsQueryExpr()  
             where num % 2 == 0  
             select num * num).Sum();
```

## compiles to

```
int sum = 0;  
for (int index = 0; index < nums.Length; index++)  
{  
    int num = nums[index];  
    if (num % 2 == 0)  
        sum += num * num;  
}
```

# Disadvantages

- Runtime compilation
  - Overhead (mitigated by caching)
  - Emitting IL not cross-platform (e.g. security restrictions in cloud, mobile)
  - Access to private fields/methods?
- Problematic F# support
- New operations => compiler changes



Should become a Roslyn compile time plugin in future

# Clash of the Lamdas

ICOOOLPS'14


Aggelos Biboudis (@biboudis)

Nick Palladinos (@NickPalladinos)


Yannis Smaragdakis

# Performance Benchmarks


Sum (windows)



## Sum (linux)

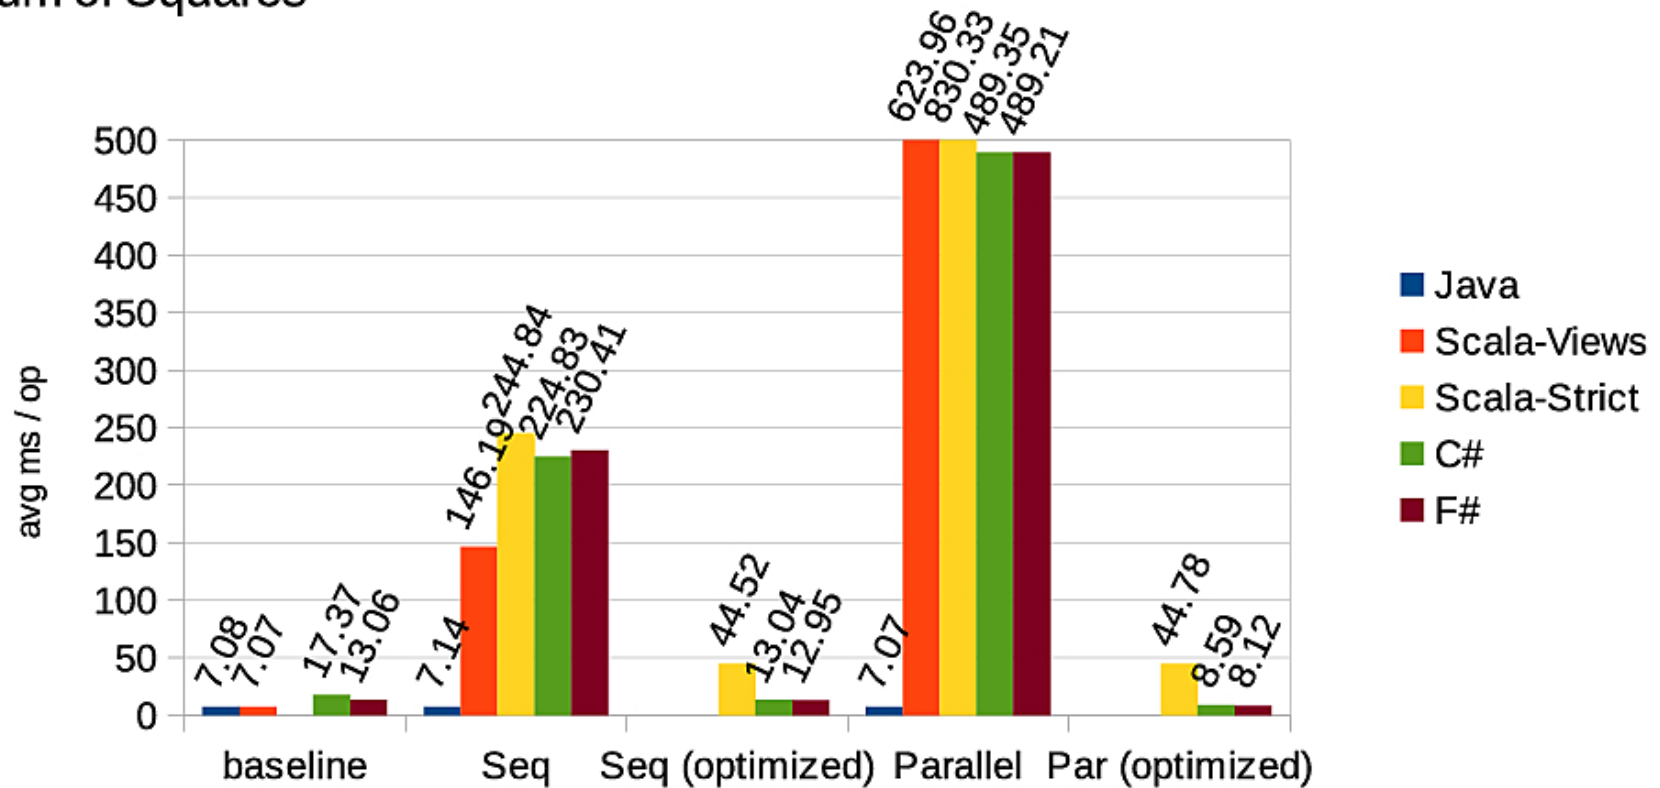


## Sum of squares (windows)




# Sum of squares (linux)


## Sum of Squares



## Sum of even squares (windows)




## Sum of even squares (linux)






# Cartesian product (windows)



# Cartesian product (linux)



Java 8 very fast

LinqOptimizer improving F#/C# performance

# What makes Java 8 faster?

# Streams!

# Typical Pipeline Pattern

```
1: source |> inter |> inter |> inter |> terminal
```

- inter : intermediate (lazy) operations, e.g. map, filter
- terminal : produces result or side-effects, e.g. reduce, iter

# Seq example

```
1: let data = [| 1..10000000 |] |> Array.map int64
2: data
3: |> Seq.filter (fun i -> i % 2L = 0L) //lazy
4: |> Seq.map (fun i -> i + 1L) //lazy
5: |> Seq.sum //eager, forcing evaluation
```



# Seq is pulling

```
1: let data = [| 1..100000000 |] |> Array.map int64
2: data
3: |> Seq.filter (fun i -> i % 2L = 0L) //lazy inter
4: |> Seq.map (fun i -> i + 1L) //lazy inter
5: |> Seq.sum //eager terminal, forcing evaluation
```

The terminal is pulling data from the pipeline via  
IEnumerator.Current and IEnumerator.MoveNext()

# With Streams

```
1: let data = [| 1..10000000 |] |> Array.map int64
2: Stream.ofArray data //source
3: |> Stream.filter (fun i -> i % 2L = 0L) //lazy
4: |> Stream.map (fun i -> i + 1L) //lazy
5: |> Stream.sum //eager, forcing evaluation
```

Streams are pushing!

# Streams are pushing

```
1: Stream.ofArray data //source
2: |> Stream.filter (fun i -> i % 2L = 0L) //lazy
3: |> Stream.map (fun i -> i + 1L) //lazy
4: |> Stream.sum //eager, forcing evaluation
```

The source is pushing data down the pipeline.

How does it work?

# Starting from Seq.iter

```
1: Seq.iter : ('T -> unit) -> seq<'T> -> unit
```

# Flip the args

```
1: seq<'T> -> ('T -> unit) -> unit
```

# Stream!

```
1: type Stream<'T> = ('T -> unit) -> unit
```



# Continuation passing style!

Let's make us some (simple) Streams!

# Simple Streams

```
1: type Stream = ('T -> unit) -> unit
```

Can do map, filter, fold, iter

# When to stop pushing?

```
1: type Stream = ('T -> unit) -> unit
```

Stopping push required for e.g.

```
1: Stream.takeWhile : ('T -> bool) -> Stream<'T> -> Stream<'T>
```

# Stopping push Change

```
1: type Stream = ('T -> unit) -> unit
```

to

```
1: type Stream = ('T -> bool) -> unit
```

# What about zip?

```
1: Stream.zip : Stream<'T> -> Stream<'S> -> Stream<'T * 'S>
```

Zip needs to synchronise the flow of values.

Zip needs to pull!

# Streams can push and pull

```
1: // ('T -> bool) is the composed continuation with 'T for the current value
2: // and bool is a flag for early termination
3: // (unit -> unit) is a function for bulk processing
4: // (unit -> bool) is a function for on-demand processing
5:
6: /// Represents a Stream of values.
7: type Stream<'T> = Stream of (('T -> bool) -> (unit -> unit) * (unit -> bool))
```

# The Streams library

Implements a rich set of operations



More examples

# Parallel Streams

```
1: let data = [| 1..10000000 |] |> Array.map int64
2: data
3: |> ParStream.ofArray
4: |> ParStream.filter (fun x -> x % 2L = 0L)
5: |> ParStream.map (fun x -> x + 1L)
6: |> ParStream.sum
```

# Cloud Streams!

Example: a word count

# Streams are lightweight and powerful

In sequential, parallel and distributed flavors.

# The holy grail is in reach

We can write functional pipelines with the performance of imperative code.

*Stream fusion: from lists to streams to nothing at all, Duncan Coutts, Roman Leshchinskiy, and Don Stewart, ICFP '07*

# Almost

Depends on the compiler's ability to inline.

Inlining continuations = stream fusion

# Stream operations are non-recursive

In principal, can be always fused (in-lined).

Not always done by F# compiler.



# Experiments with MLton

by @biboudis

<https://github.com/biboudis/sml-streams>

MLton appears to always be fusing.

Can we make the F# compiler smarter?

Questions?